

Build End-to-End ML pipeline for Warfarin Dosing Prediction

```
# Reading excel file
import pandas as pd
import numpy as np
data_frame = pd.read_excel('FinalData.xls')
data_frame.head()
```



| | Gender | Race (Reported) | Age | Height (cm) | Weight (kg) | Diabetes | Simvastatin (Zocor) | Amiodarone (Cordarone) |
|---|--------|--------------------|------------|----------------|----------------|----------|------------------------|---------------------------|
| 0 | male | White | 60 - 69 | 193.040 | 115.7 | NaN | 0.0 | 0.0 |
| 1 | female | White | 50 - 59 | 176.530 | 144.2 | NaN | 0.0 | 0.0 |
| 2 | female | White | 40 - 49 | 162.560 | 77.1 | NaN | 0.0 | 0.0 |
| 3 | male | White | 60 - 69 | 182.245 | 90.7 | NaN | 0.0 | 0.0 |
| 4 | male | White | 50 - 59 | 167.640 | 72.6 | NaN | 0.0 | 0.0 |

```
from google.colab import drive
drive.mount('/content/drive')
```



Mounted at /content/drive

```
# Checking columns of the dataset
data_frame.columns
```

```
Index(['Gender', 'Race (Reported)', 'Age', 'Height (cm)', 'Weight (kg)',
      'Diabetes', 'Simvastatin (Zocor)', 'Amiodarone (Cordarone)',
      'Target INR', 'INR on Reported Therapeutic Dose of Warfarin',
      'Cyp2C9 genotypes',
      'VKORC1 genotype: -1639 G>A (3673); chr16:31015190; rs9923231; C/T',
      'Therapeutic Dose of Warfarin'],
      dtype='object')
```

```
# Check missing values in the dataset
data_frame.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5700 entries, 0 to 5699
Data columns (total 13 columns):
#   Column                                                                 Non-N
---  -
0   Gender                                                                5696
1   Race (Reported)                                                       5194
2   Age                                                                    5658
3   Height (cm)                                                            4554
4   Weight (kg)                                                            5413
5   Diabetes                                                                3283
6   Simvastatin (Zocor)                                                    3861
7   Amiodarone (Cordarone)                                                 4182
8   Target INR                                                             1259
9   INR on Reported Therapeutic Dose of Warfarin                         4968
10  Cyp2C9 genotypes                                                       5567
11  VKORC1 genotype: -1639 G>A (3673); chr16:31015190; rs9923231; C/T  4046
12  Therapeutic Dose of Warfarin                                           5528
dtypes: float64(8), object(5)
memory usage: 579.0+ KB
```

```
# Describing the dataset
data_frame.describe()
```



| | Height (cm) | Weight (kg) | Diabetes | Simvastatin (Zocor) | Amiodarone (Cordarone) | Target INR |
|-------|----------------|----------------|-------------|------------------------|---------------------------|---------------|
| count | 4554.000000 | 5413.000000 | 3283.000000 | 3861.000000 | 4182.000000 | 1259.000000 |
| mean | 168.047778 | 77.852569 | 0.187024 | 0.146335 | 0.066236 | 2.538324 |
| std | 10.845992 | 21.859764 | 0.389990 | 0.353488 | 0.248724 | 0.198140 |
| min | 124.968000 | 30.000000 | 0.000000 | 0.000000 | 0.000000 | 1.300000 |
| 25% | 160.020000 | 62.000000 | 0.000000 | 0.000000 | 0.000000 | 2.500000 |
| 50% | 167.894000 | 75.000000 | 0.000000 | 0.000000 | 0.000000 | 2.500000 |
| 75% | 176.022000 | 90.000000 | 0.000000 | 0.000000 | 0.000000 | 2.500000 |

```
# Checking for NA Values in the dataset
data_frame.isna()
```



| | Gender | Race (Reported) | Age | Height (cm) | Weight (kg) | Diabetes | Simvastatin (Zocor) | Amiodarone (Cordarone) |
|------------------------|--------|--------------------|-------|----------------|----------------|----------|------------------------|---------------------------|
| 0 | False | False | False | False | False | True | False | False |
| 1 | False | False | False | False | False | True | False | False |
| 2 | False | False | False | False | False | True | False | False |
| 3 | False | False | False | False | False | True | False | False |
| 4 | False | False | False | False | False | True | False | False |
| ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 5695 | False | False | False | False | False | False | False | False |
| 5696 | False | False | False | False | False | False | False | False |
| 5697 | False | False | False | False | False | False | False | False |
| 5698 | False | False | False | False | False | False | False | False |
| 5699 | False | False | False | False | False | False | False | False |
| 5700 rows × 13 columns | | | | | | | | |

```
# Finding Total Missing Values
data_frame.isna().sum()
```

```
⇒ Gender 4
Race (Reported) 506
Age 42
Height (cm) 1146
Weight (kg) 287
Diabetes 2417
Simvastatin (Zocor) 1839
Amiodarone (Cordarone) 1518
Target INR 4441
INR on Reported Therapeutic Dose of Warfarin 732
Cyp2C9 genotypes 133
VKORC1 genotype: -1639 G>A (3673); chr16:31015190; rs9923231; C/T 1654
Therapeutic Dose of Warfarin 172
dtype: int64
```

```
# Total sum of missing values
data_frame.isna().sum().sum()
```

```
⇒ 14891
```

```
# Replacing missing values with mode
```

```
mode_fillers = {
    'Gender': data_frame['Gender'].mode()[0],
    'Race (Reported)': data_frame['Race (Reported)'].mode()[0],
    'Age': data_frame['Age'].mode()[0],
    'Cyp2C9 genotypes': data_frame['Cyp2C9 genotypes'].mode()[0],
    'VKORC1 genotype: -1639 G>A (3673); chr16:31015190; rs9923231; C/T': data_fra
}
```

```
data_frame.fillna(value=mode_fillers, inplace=True)
```

```
# Replacing missing values with mean
mean_fillers = {
    'Height (cm)': data_frame['Height (cm)'].mean(),
    'Weight (kg)': data_frame['Weight (kg)'].mean(),
    'Diabetes': data_frame['Diabetes'].mean(),
    'Simvastatin (Zocor)': data_frame['Simvastatin (Zocor)'].mean(),
    'Amiodarone (Cordarone)': data_frame['Amiodarone (Cordarone)'].mean(),
    'Target INR': data_frame['Target INR'].mean(),
    'INR on Reported Therapeutic Dose of Warfarin': data_frame['INR on Reported T
    'Therapeutic Dose of Warfarin': data_frame['Therapeutic Dose of Warfarin'].me
}

data_frame.fillna(value=mean_fillers, inplace=True)
```

```
# checking for missing values
data_frame.isna()
```



| | Gender | Race (Reported) | Age | Height (cm) | Weight (kg) | Diabetes | Simvastatin (Zocor) | Amiodarone (Cordarone) |
|------------------------|--------|--------------------|-------|----------------|----------------|----------|------------------------|---------------------------|
| 0 | False | False | False | False | False | False | False | False |
| 1 | False | False | False | False | False | False | False | False |
| 2 | False | False | False | False | False | False | False | False |
| 3 | False | False | False | False | False | False | False | False |
| 4 | False | False | False | False | False | False | False | False |
| ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 5695 | False | False | False | False | False | False | False | False |
| 5696 | False | False | False | False | False | False | False | False |
| 5697 | False | False | False | False | False | False | False | False |
| 5698 | False | False | False | False | False | False | False | False |
| 5699 | False | False | False | False | False | False | False | False |
| 5700 rows × 13 columns | | | | | | | | |

```
data_frame.isna().sum()
```

```

⇒ Gender 0
  Race (Reported) 0
  Age 0
  Height (cm) 0
  Weight (kg) 0
  Diabetes 0
  Simvastatin (Zocor) 0
  Amiodarone (Cordarone) 0
  Target INR 0
  INR on Reported Therapeutic Dose of Warfarin 0
  Cyp2C9 genotypes 0
  VKORC1 genotype: -1639 G>A (3673); chr16:31015190; rs9923231; C/T 0
  Therapeutic Dose of Warfarin 0
  dtype: int64

```

```

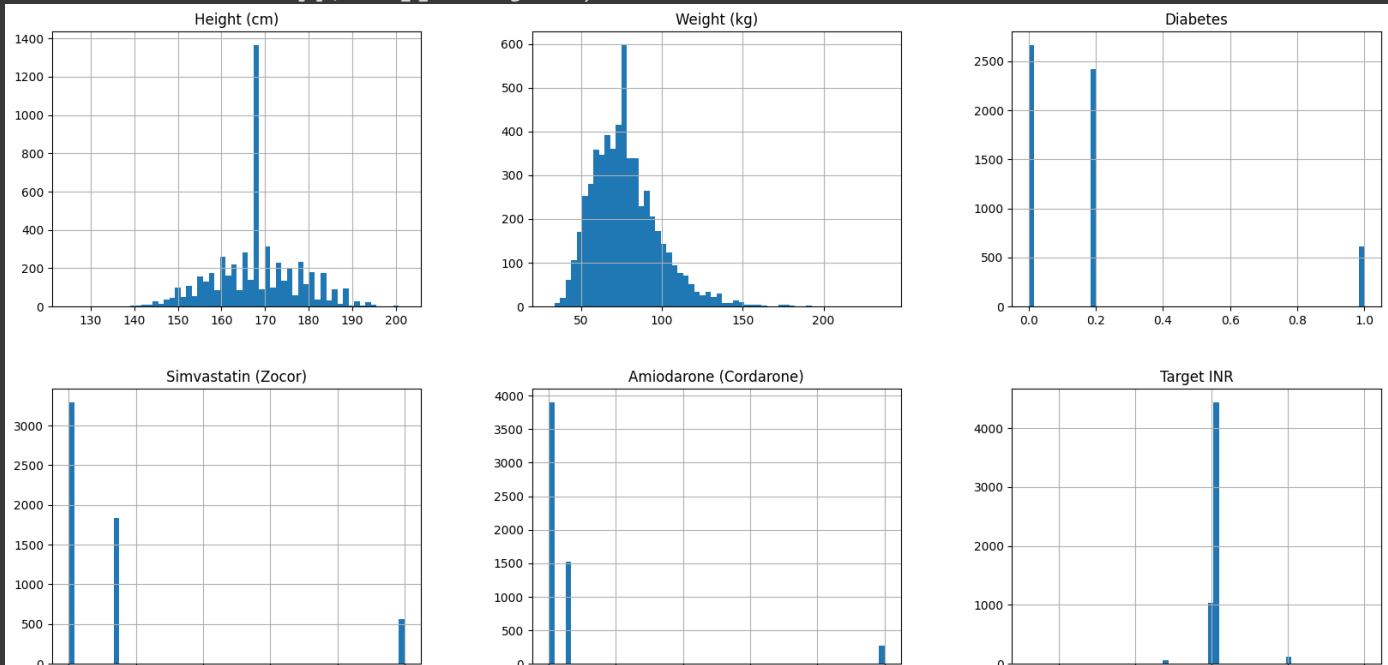
# visualization of dataset
data_frame.hist(bins=60, figsize=(20,15))


```

```

⇒ array([[<Axes: title={'center': 'Height (cm)'}>,
  <Axes: title={'center': 'Weight (kg)'}>,
  <Axes: title={'center': 'Diabetes'}>],
  [<Axes: title={'center': 'Simvastatin (Zocor)'}>,
  <Axes: title={'center': 'Amiodarone (Cordarone)'}>,
  <Axes: title={'center': 'Target INR'}>],
  [<Axes: title={'center': 'INR on Reported Therapeutic Dose of
Warfarin'}>,
  <Axes: title={'center': 'Therapeutic Dose of Warfarin'}>,
  <Axes: >]], dtype=object)

```





Data visualization

```
# Creating labels for model
df_target = data_frame[['Therapeutic Dose of Warfarin']]
df_features = data_frame.drop(columns=['Therapeutic Dose of Warfarin'])
```



```
# Features without Target column
df_features.columns
```

```
Index(['Gender', 'Race (Reported)', 'Age', 'Height (cm)', 'Weight (kg)',
      'Diabetes', 'Simvastatin (Zocor)', 'Amiodarone (Cordarone)',
      'Target INR', 'INR on Reported Therapeutic Dose of Warfarin',
      'Cyp2C9 genotypes',
      'VKORC1 genotype: -1639 G>A (3673); chr16:31015190; rs9923231; C/T'],
      dtype='object')
```

```
# Target column
df_target.columns
```

```
Index(['Therapeutic Dose of Warfarin'], dtype='object')
```

```
# Splitting the data into train and test
from sklearn.model_selection import train_test_split
split_data = train_test_split(df_features, df_target, test_size=0.2, random_state=42)
X_train, X_test, y_train, y_test = split_data
X_train.head()
```

```
Index(['Gender', 'Race (Reported)', 'Age', 'Height (cm)', 'Weight (kg)', 'Diabetes', 'Simvastatin (Zocor)', 'Amiodarone (Cordarone)', 'Target INR', 'INR on Reported Therapeutic Dose of Warfarin', 'Cyp2C9 genotypes', 'VKORC1 genotype: -1639 G>A (3673); chr16:31015190; rs9923231; C/T'], dtype='object')
```

| | Gender | Race (Reported) | Age | Height (cm) | Weight (kg) | Diabetes | Simvastatin (Zocor) | Amiodarone (Cordarone) | Target INR | INR on Reported Therapeutic Dose of Warfarin | Cyp2C9 genotypes | VKORC1 genotype: -1639 G>A (3673); chr16:31015190; rs9923231; C/T |
|------|--------|-----------------|---------|-------------|-------------|----------|---------------------|------------------------|------------|--|------------------|---|
| 5437 | female | White | 70 - 79 | 160.020000 | 60.91 | 0.000000 | 0.000000 | 0.000000 | 0.000000 | 0.000000 | 0.000000 | 0.000000 |
| 2979 | male | Korean | 20 - 29 | 178.000000 | 82.00 | 0.000000 | 0.000000 | 0.000000 | 0.000000 | 0.000000 | 0.000000 | 0.000000 |
| 4743 | female | Malay | 60 - 69 | 168.047778 | 72.00 | 0.187024 | 0.146335 | 0.000000 | 0.000000 | 0.000000 | 0.000000 | 0.000000 |
| 2668 | male | White | 70 - 79 | 175.260000 | 85.30 | 0.000000 | 0.000000 | 0.000000 | 0.000000 | 0.000000 | 0.000000 | 0.000000 |
| 3264 | male | Caucasian | 80 - 89 | 168.047778 | 59.00 | 0.000000 | 0.146335 | 0.000000 | 0.000000 | 0.000000 | 0.000000 | 0.000000 |

```
y_test.head()
```



| Therapeutic Dose of Warfarin | |
|------------------------------|-------|
| 1477 | 17.50 |
| 5514 | 37.52 |
| 3243 | 13.72 |
| 5183 | 21.00 |
| 5107 | 14.00 |

```
# Transformation of categorical values to binary of train set
from sklearn import preprocessing
le = preprocessing.LabelEncoder()
X_train['Race (Reported)'] = le.fit_transform(X_train['Race (Reported)'])
X_train['Gender'] = le.fit_transform(X_train['Gender'])
X_train['Age'] = le.fit_transform(X_train['Age'])
X_train['Cyp2C9 genotypes'] = le.fit_transform(X_train['Cyp2C9 genotypes'])
X_train['VKORC1 genotype: -1639 G>A (3673); chr16:31015190; rs9923231; C/T'] = le
```

```
# MinMax normalization of train set
from sklearn.preprocessing import MinMaxScaler
scaler = MinMaxScaler() ## define the transformer
scaler.fit(X_train) ## call .fit() method to calculate the min and max value for
x_train_normalized = scaler.transform(X_train)
x_train_normalized=pd.DataFrame(x_train_normalized)
x_train_normalized
```



| | 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | |
|------|-----|------|-------|----------|----------|----------|----------|----------|----------|--------|
| 0 | 0.0 | 0.95 | 0.750 | 0.455032 | 0.148820 | 0.000000 | 0.000000 | 0.000000 | 0.450471 | 0.2641 |
| 1 | 1.0 | 0.65 | 0.125 | 0.688441 | 0.250361 | 0.000000 | 0.000000 | 0.000000 | 0.450471 | 0.3150 |
| 2 | 0.0 | 0.70 | 0.625 | 0.559245 | 0.202215 | 0.187024 | 0.146335 | 0.066236 | 0.450471 | 0.2490 |
| 3 | 1.0 | 0.95 | 0.750 | 0.652872 | 0.266249 | 0.000000 | 0.000000 | 0.000000 | 0.450471 | 0.3207 |
| 4 | 1.0 | 0.30 | 0.875 | 0.559245 | 0.139624 | 0.000000 | 0.146335 | 0.000000 | 0.450471 | 0.2951 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 4555 | 1.0 | 0.60 | 0.375 | 0.596817 | 0.235917 | 0.187024 | 0.146335 | 0.066236 | 0.450471 | 0.2951 |
| 4556 | 0.0 | 0.10 | 0.625 | 0.586925 | 0.474723 | 1.000000 | 1.000000 | 0.000000 | 0.428571 | 0.3584 |
| 4557 | 0.0 | 0.60 | 0.625 | 0.375896 | 0.144439 | 0.187024 | 0.146335 | 0.066236 | 0.450471 | 0.1075 |
| 4558 | 0.0 | 0.95 | 0.375 | 0.455032 | 0.144439 | 0.000000 | 0.000000 | 0.000000 | 0.450471 | 0.3773 |
| 4559 | 0.0 | 0.60 | 0.625 | 0.428653 | 0.144439 | 0.187024 | 0.146335 | 0.066236 | 0.450471 | 0.4132 |

4560 rows × 12 columns

```
# Transformation of categorical values to binary of test set
from sklearn import preprocessing
le = preprocessing.LabelEncoder()
X_test['Race (Reported)'] = le.fit_transform(X_test['Race (Reported)'])
X_test['Gender'] = le.fit_transform(X_test['Gender'])
X_test['Age'] = le.fit_transform(X_test['Age'])
X_test['Cyp2C9 genotypes'] = le.fit_transform(X_test['Cyp2C9 genotypes'])
X_test['VKORC1 genotype: -1639 G>A (3673); chr16:31015190; rs9923231; C/T'] = le.
```

```
# MinMax normilization of train set
from sklearn.preprocessing import MinMaxScaler

scaler = MinMaxScaler() ## define the transformer
scaler.fit(X_test) ## call .fit() method to calculate the min and max value for e
x_test_normalized = scaler.transform(X_test)
x_test_normalized=pd.DataFrame(x_test_normalized)
x_test_normalized
```



| | 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 |
|------|-----|--------|----------|----------|----------|----------|----------|----------|----------|
| 0 | 0.0 | 0.4375 | 0.142857 | 0.332917 | 0.241880 | 0.000000 | 0.000000 | 0.000000 | 0.562875 |
| 1 | 1.0 | 0.2500 | 0.714286 | 0.833333 | 0.339254 | 0.000000 | 0.000000 | 0.000000 | 0.562875 |
| 2 | 0.0 | 0.3125 | 0.571429 | 0.465023 | 0.297167 | 0.000000 | 0.146335 | 1.000000 | 0.562875 |
| 3 | 1.0 | 1.0000 | 0.285714 | 0.546250 | 0.290256 | 0.000000 | 0.000000 | 1.000000 | 0.562875 |
| 4 | 0.0 | 1.0000 | 0.857143 | 0.300000 | 0.345543 | 0.000000 | 0.000000 | 0.000000 | 0.562875 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 1135 | 1.0 | 0.6875 | 0.285714 | 0.366667 | 0.172771 | 0.187024 | 0.146335 | 0.066236 | 0.562875 |
| 1136 | 1.0 | 0.3125 | 0.142857 | 0.465023 | 0.200415 | 0.000000 | 0.146335 | 0.000000 | 0.562875 |
| 1137 | 1.0 | 1.0000 | 0.285714 | 0.637500 | 0.373186 | 0.000000 | 1.000000 | 0.000000 | 0.545455 |
| 1138 | 1.0 | 0.6875 | 0.571429 | 0.445833 | 0.124395 | 0.187024 | 0.146335 | 0.066236 | 0.562875 |
| 1139 | 1.0 | 0.6875 | 0.571429 | 0.662500 | 0.262612 | 0.187024 | 0.146335 | 0.066236 | 0.562875 |

1140 rows × 12 columns

```
# Binary classification dataset by cutting the target values into two categories
y_train[y_train<=30] = 0
y_train[y_train>30] = 1
```

y_train



Therapeutic Dose of Warfarin

| | |
|------|-----|
| 5437 | 0.0 |
| 2979 | 1.0 |
| 4743 | 0.0 |
| 2668 | 0.0 |
| 3264 | 1.0 |
| ... | ... |
| 1180 | 0.0 |
| 3441 | 1.0 |
| 1344 | 0.0 |
| 4623 | 0.0 |
| 1289 | 0.0 |

4560 rows × 1 columns

```
# Display unique values of the 'Therapeutic Dose of Warfarin' column in train set  
print(y_train["Therapeutic Dose of Warfarin"].unique())
```



```
[0. 1.]
```

```
y_test[y_test<=30] = 0  
y_test[y_test>30] = 1
```

y_test



Therapeutic Dose of Warfarin

| | |
|------|-----|
| 1477 | 0.0 |
| 5514 | 1.0 |
| 3243 | 0.0 |
| 5183 | 0.0 |
| 5107 | 0.0 |
| ... | ... |
| 924 | 0.0 |
| 3317 | 0.0 |
| 1846 | 0.0 |
| 1316 | 0.0 |
| 1064 | 0.0 |

1140 rows × 1 columns

```
# Display unique values of the 'Therapeutic Dose of Warfarin' column in test set  
print(y_test["Therapeutic Dose of Warfarin"].unique())
```



[0. 1.]

Model training

```
# Decision Tree
from sklearn.tree import DecisionTreeClassifier

# Define the decision tree model
decision_tree_model = DecisionTreeClassifier(max_depth=3)

# Fit the model on the normalized training data
decision_tree_model.fit(x_train_normalized, y_train)

# Predict using the trained model on the normalized test data
decision_tree_model.predict(x_test_normalized)
```

⇒ array([1., 1., 0., ..., 1., 0., 1.])

```
# calculating score
decision_tree_model.score(x_train_normalized,y_train)
```

⇒ 0.7032894736842106

```
# Decision tree performance on test data
from sklearn.metrics import accuracy_score, precision_score, recall_score, roc_auc_score

# Predict using the decision tree model on the normalized test data
decision_tree_model_pred = decision_tree_model.predict(x_test_normalized)

# Calculate accuracy
decision_tree_model_acc = accuracy_score(y_test, decision_tree_model_pred)

# Calculate precision
decision_tree_model_prec = precision_score(y_test, decision_tree_model_pred)

# Calculate recall
decision_tree_model_recall = recall_score(y_test, decision_tree_model_pred)

# Calculate ROC AUC score
decision_tree_model_roc = roc_auc_score(y_test, decision_tree_model_pred)

# Calculate F1 score
decision_tree_model_f1 = f1_score(y_test, decision_tree_model_pred)

# Print performance metrics
print(decision_tree_model_acc)
print(decision_tree_model_prec)
print(decision_tree_model_recall)
print(decision_tree_model_roc)
print(decision_tree_model_f1)
```

```
0.637719298245614
0.5602165087956699
0.8247011952191236
0.657648403252195
0.6672038678485093
```

```
!pip install --upgrade scikit-learn
```

```
Requirement already satisfied: scikit-learn in /usr/local/lib/python3.10/dist/
Requirement already satisfied: numpy>=1.19.5 in /usr/local/lib/python3.10/dist/
Requirement already satisfied: scipy>=1.6.0 in /usr/local/lib/python3.10/dist/
Requirement already satisfied: joblib>=1.2.0 in /usr/local/lib/python3.10/dist/
Requirement already satisfied: threadpoolctl>=2.0.0 in /usr/local/lib/python3
```

```
from sklearn.metrics import roc_curve, auc
```




```
import matplotlib.pyplot as plt

# Assuming decision_tree_model is your classifier and X_test, y_test are your test data
y_score = decision_tree_model.predict_proba(X_test)[:,1]
fpr, tpr, _ = roc_curve(y_test, y_score)
roc_auc = auc(fpr, tpr)

plt.figure()
plt.plot(fpr, tpr, color='darkorange', lw=2, label='ROC curve (area = %0.2f)' % roc_auc)
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend(loc="lower right")
plt.show()
```



```
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:486: UserWarning: X h
warnings.warn(
```



```
# visualization of decision tree
import matplotlib.pyplot as plt
from sklearn import tree

fig = plt.figure(figsize=(20, 10))
tree.plot_tree(decision_tree_model)
```

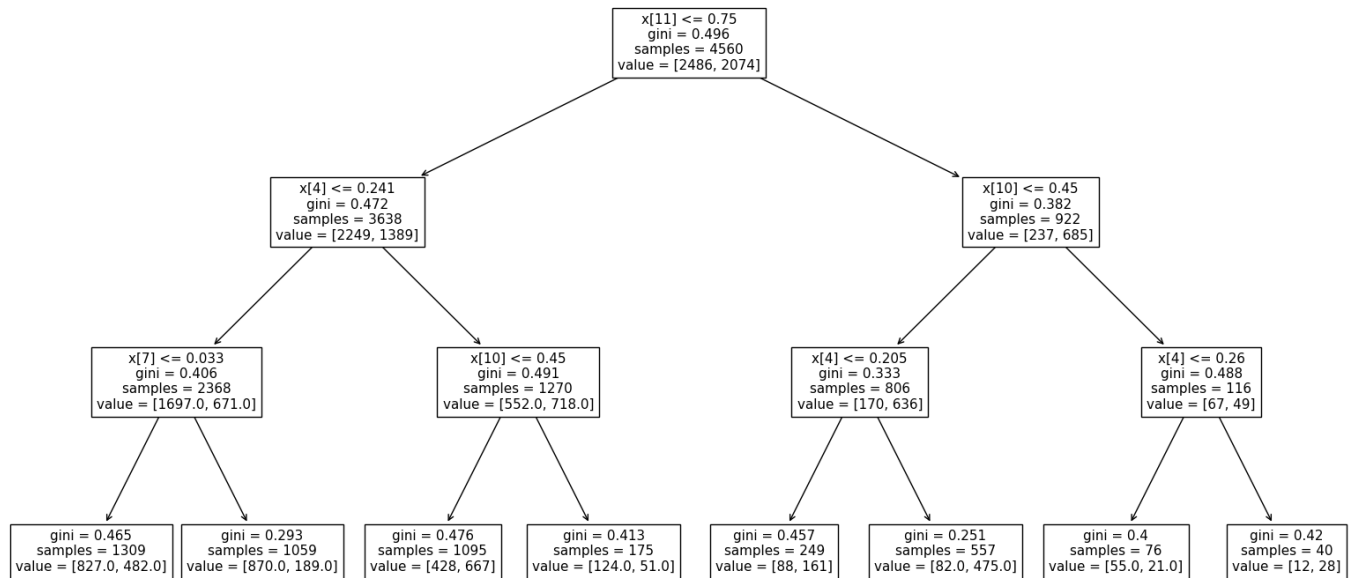


```
[Text(0.5, 0.875, 'x[11] <= 0.75\ngini = 0.496\nsamples = 4560\nvalue =
[2486, 2074]'),
Text(0.25, 0.625, 'x[4] <= 0.241\ngini = 0.472\nsamples = 3638\nvalue =
[2249, 1389]'),
Text(0.125, 0.375, 'x[7] <= 0.033\ngini = 0.406\nsamples = 2368\nvalue =
[1697.0, 671.0]'),
Text(0.0625, 0.125, 'gini = 0.465\nsamples = 1309\nvalue = [827.0, 482.0]'),
Text(0.1875, 0.125, 'gini = 0.293\nsamples = 1059\nvalue = [870.0, 189.0]'),
Text(0.375, 0.375, 'x[10] <= 0.45\ngini = 0.491\nsamples = 1270\nvalue =
```

```

[552.0, 718.0]'),
Text(0.3125, 0.125, 'gini = 0.476\nsamples = 1095\nvalue = [428, 667]'),
Text(0.4375, 0.125, 'gini = 0.413\nsamples = 175\nvalue = [124.0, 51.0]'),
Text(0.75, 0.625, 'x[10] <= 0.45\ngini = 0.382\nsamples = 922\nvalue = [237, 685]'),
Text(0.625, 0.375, 'x[4] <= 0.205\ngini = 0.333\nsamples = 806\nvalue = [170, 636]'),
Text(0.5625, 0.125, 'gini = 0.457\nsamples = 249\nvalue = [88, 161]'),
Text(0.6875, 0.125, 'gini = 0.251\nsamples = 557\nvalue = [82.0, 475.0]'),
Text(0.875, 0.375, 'x[4] <= 0.26\ngini = 0.488\nsamples = 116\nvalue = [67, 49]'),
Text(0.8125, 0.125, 'gini = 0.4\nsamples = 76\nvalue = [55.0, 21.0]'),
Text(0.9375, 0.125, 'gini = 0.42\nsamples = 40\nvalue = [12, 28]')]

```



```
import joblib
model_file_path = "/content/drive/My Drive/decision_tree_model.pkl"
#joblib.dump(decision_tree_model, model_file_path)
```

```
# Logistic regression
from sklearn.linear_model import LogisticRegression

# Create a logistic regression model
logistic_model = LogisticRegression(penalty='l2', C=1, random_state=0)

# Fit the model to the normalized training data
logistic_model.fit(x_train_normalized, y_train)

# Make predictions on the normalized test data
logistic_predictions = logistic_model.predict(x_test_normalized)

# Print the predictions
print(logistic_predictions)
```

```
⇒ [0. 1. 0. ... 1. 0. 0.]
/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:1300: Dat
y = column_or_1d(y, warn=True)
```

```
logistic_model.score(x_train_normalized,y_train)
```

```
⇒ 0.7247807017543859
```

```
# Logistic regression – performance on the test data
from sklearn.metrics import accuracy_score, precision_score, recall_score, roc_auc_score

# Predicting the target values for the test data using the logistic regression model
logistic_pred = logistic_model.predict(x_test_normalized)

# Calculating the accuracy score
logistic_acc = accuracy_score(y_test, logistic_pred)

# Calculating the precision score
logistic_prec = precision_score(y_test, logistic_pred)

# Calculating the recall score
logistic_recall = recall_score(y_test, logistic_pred)

# Calculating the ROC AUC score
logistic_roc = roc_auc_score(y_test, logistic_pred)

# Calculating the F1 score
logistic_f1 = f1_score(y_test, logistic_pred)

# Printing the calculated performance metrics
print(logistic_acc)
print(logistic_prec)
print(logistic_recall)
print(logistic_roc)
print(logistic_f1)
```

```
0.7035087719298245
0.6238670694864048
0.8227091633466136
0.7162135158425857
0.7096219931271478
```

```
from sklearn.metrics import roc_curve, auc
import matplotlib.pyplot as plt


# Assuming logistic_model is your classifier and X_test, y_test are your test data
y_score = logistic_model.predict_proba(X_test)[:,1]
fpr, tpr, _ = roc_curve(y_test, y_score)
roc_auc = auc(fpr, tpr)

plt.figure()
```

```
plt.plot(fpr, tpr, color='darkorange', lw=2, label='Logistic Regression (area = %s)' % area)
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend(loc="lower right")
plt.show()
```



/usr/local/lib/python3.10/dist-packages/sklearn/base.py:486: UserWarning: X has feature names, but LogisticRegression does not allow them. Please use LogisticRegressionCV instead.
warnings.warn(msg, UserWarning)



```
import joblib
model_file_path = "/content/drive/My Drive/logistic_model.pkl"
joblib.dump(logistic_model, model_file_path)
```



['/content/drive/My Drive/logistic_model.pkl']

```
# Support Vector Machine
from sklearn.svm import SVC

# Define SVM model with C=1
svm_model = SVC(C=1)

# Fit SVM model to the normalized training data
svm_model.fit(x_train_normalized, y_train)

# Predict target values for the normalized test data
svm_model.predict(x_test_normalized)

# Compute and print the accuracy score of the SVM model on the training data
svm_model.score(x_train_normalized, y_train)
```

➞ /usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:1300: Dat
y = column_or_1d(y, warn=True)
0.7449561403508772

```
import joblib
model_file_path = "/content/drive/My Drive/svm_model.pkl"
joblib.dump(svm_model, model_file_path)
```

➞ ['/content/drive/My Drive/svm_model.pkl']

```
# linear kernel
svm_model_linear = SVC(kernel="linear", degree=3, C=5)

# Fit SVM model with linear kernel to the normalized training data
svm_model_linear.fit(x_train_normalized, y_train)

# Predict target values for the normalized test data
svm_model_linear.predict(x_test_normalized)

# Compute and print the accuracy score of the SVM model with linear kernel on the
svm_model_linear.score(x_train_normalized, y_train)
```

➞ /usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:1300: Dat
y = column_or_1d(y, warn=True)
0.7243421052631579

```
import joblib
model_file_path = "/content/drive/My Drive/svm_model_linear.pkl"
joblib.dump(svm_model_linear, model_file_path)
```

➞ ['content/drive/My Drive/svm_model_linear.pkl']

```
# RBF kernel SVM model
svm_model_rbf = SVC(kernel="rbf", C=5)

# Train the SVM model with RBF kernel using the normalized training data
svm_model_rbf.fit(x_train_normalized, y_train)

# Predict target values for the normalized test data
svm_model_rbf.predict(x_test_normalized)

# Compute and print the accuracy score of the SVM model with RBF kernel on the tr
svm_model_rbf.score(x_train_normalized, y_train)
```

➞ /usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:1300: Dat
y = column_or_1d(y, warn=True)
0.756578947368421

```
import joblib
model_file_path = "/content/drive/My Drive/svm_model_rbf.pkl"
joblib.dump(svm_model_rbf, model_file_path)
```

➞ ['content/drive/My Drive/svm_model_rbf.pkl']


```
# Sigmoid kernel SVM model with degree included
svm_model_sigmoid = SVC(kernel="sigmoid", degree=3, C=5)

# Train the SVM model with Sigmoid kernel using the normalized training data
svm_model_sigmoid.fit(x_train_normalized, y_train)

# Predict target values for the normalized test data
svm_model_sigmoid.predict(x_test_normalized)

# Compute and print the accuracy score of the SVM model with Sigmoid kernel on th
svm_model_sigmoid.score(x_train_normalized, y_train)
```

```
➞ /usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:1300: Dat
  y = column_or_1d(y, warn=True)
0.493859649122807
```

```
import joblib
model_file_path = "/content/drive/My Drive/svm_model_sigmoid.pkl"
joblib.dump(svm_model_sigmoid, model_file_path)
```

```
➞ ['content/drive/My Drive/svm_model_sigmoid.pkl']
```

```
# Polynomial kernel SVM model
svm_model_polynomial = SVC(kernel="poly", degree=3, C=5)

# Train the SVM model with Polynomial kernel using the normalized training data
svm_model_polynomial.fit(x_train_normalized, y_train)

# Predict target values for the normalized test data
svm_model_polynomial.predict(x_test_normalized)

# Compute and print the accuracy score of the SVM model with Polynomial kernel on
svm_model_polynomial.score(x_train_normalized, y_train)
```

```
➞ /usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:1300: Dat
  y = column_or_1d(y, warn=True)
0.7526315789473684
```

```
import joblib
model_file_path = "/content/drive/My Drive/svm_model_polynomial.pkl"
joblib.dump(svm_model_polynomial, model_file_path)
```

```
➞ ['/content/drive/My Drive/svm_model_polynomial.pkl']
```

```
from sklearn.metrics import accuracy_score, precision_score, recall_score, roc_auc_score

# SVM Model Evaluation
svm_model_pred = svm_model_polynomial.predict(x_test_normalized)
svm_acc = accuracy_score(y_test, svm_model_pred)
svm_prec = precision_score(y_test, svm_model_pred)
svm_recall = recall_score(y_test, svm_model_pred)
svm_roc = roc_auc_score(y_test, svm_model_pred)
svm_f1 = f1_score(y_test, svm_model_pred)
print(svm_acc)
print(svm_prec)
print(svm_recall)
print(svm_roc)
print(svm_f1)
```

```
➞ 0.718421052631579
0.661319073083779
0.7390438247011952
0.7206190910339833
0.6980244590780809
```

```
from sklearn.metrics import roc_curve, auc
import matplotlib.pyplot as plt


# Assuming svm_model_polynomial is your classifier and X_test, y_test are your test data
y_score = svm_model_polynomial.decision_function(X_test)
fpr, tpr, _ = roc_curve(y_test, y_score)
roc_auc = auc(fpr, tpr)

plt.figure()
plt.plot(fpr, tpr, color='darkorange', lw=2, label='SVC (area = %0.2f)' % roc_auc)
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
```

```
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend(loc="lower right")
plt.show()
```



/usr/local/lib/python3.10/dist-packages/sklearn/base.py:486: UserWarning: X h
warnings.warn(



```
# Importing the K Nearest-neighbors classifier from sklearn
from sklearn.neighbors import KNeighborsClassifier

# Initializing the K Nearest-neighbors model with 5 neighbors and the Minkowski d
knn_model = KNeighborsClassifier(n_neighbors=5, metric="minkowski")

# Fitting the K Nearest-neighbors model on the normalized training data
knn_model.fit(x_train_normalized, y_train)

# Predicting the target values for the normalized test data using the trained K N
knn_model.predict(x_test_normalized)

# Calculating the accuracy score of the K Nearest-neighbors model on the normaliz
knn_model.score(x_train_normalized, y_train)
```

```
➦ /usr/local/lib/python3.10/dist-packages/sklearn/neighbors/_classification.py:
    return self._fit(X, y)
0.7916666666666666
```

```
import joblib
model_file_path = "/content/drive/My Drive/knn_model.pkl"
joblib.dump(knn_model, model_file_path)
```

```
➦ ['/content/drive/My Drive/knn_model.pkl']
```

```
# KNN – performance on the test data
from sklearn.metrics import accuracy_score, precision_score, recall_score, roc_auc_score

# KNN Model Evaluation
knn_model_pred = knn_model.predict(x_test_normalized)
knn_acc = accuracy_score(y_test, knn_model_pred)
knn_prec = precision_score(y_test, knn_model_pred)
knn_recall = recall_score(y_test, knn_model_pred)
knn_roc = roc_auc_score(y_test, knn_model_pred)
knn_f1 = f1_score(y_test, knn_model_pred)

print(knn_acc)
print(knn_prec)
print(knn_recall)
print(knn_roc)
print(knn_f1)
```

```
⇒ 0.7149122807017544
   0.6654205607476635
   0.7091633466135459
   0.714299541645331
   0.686595949855352
```

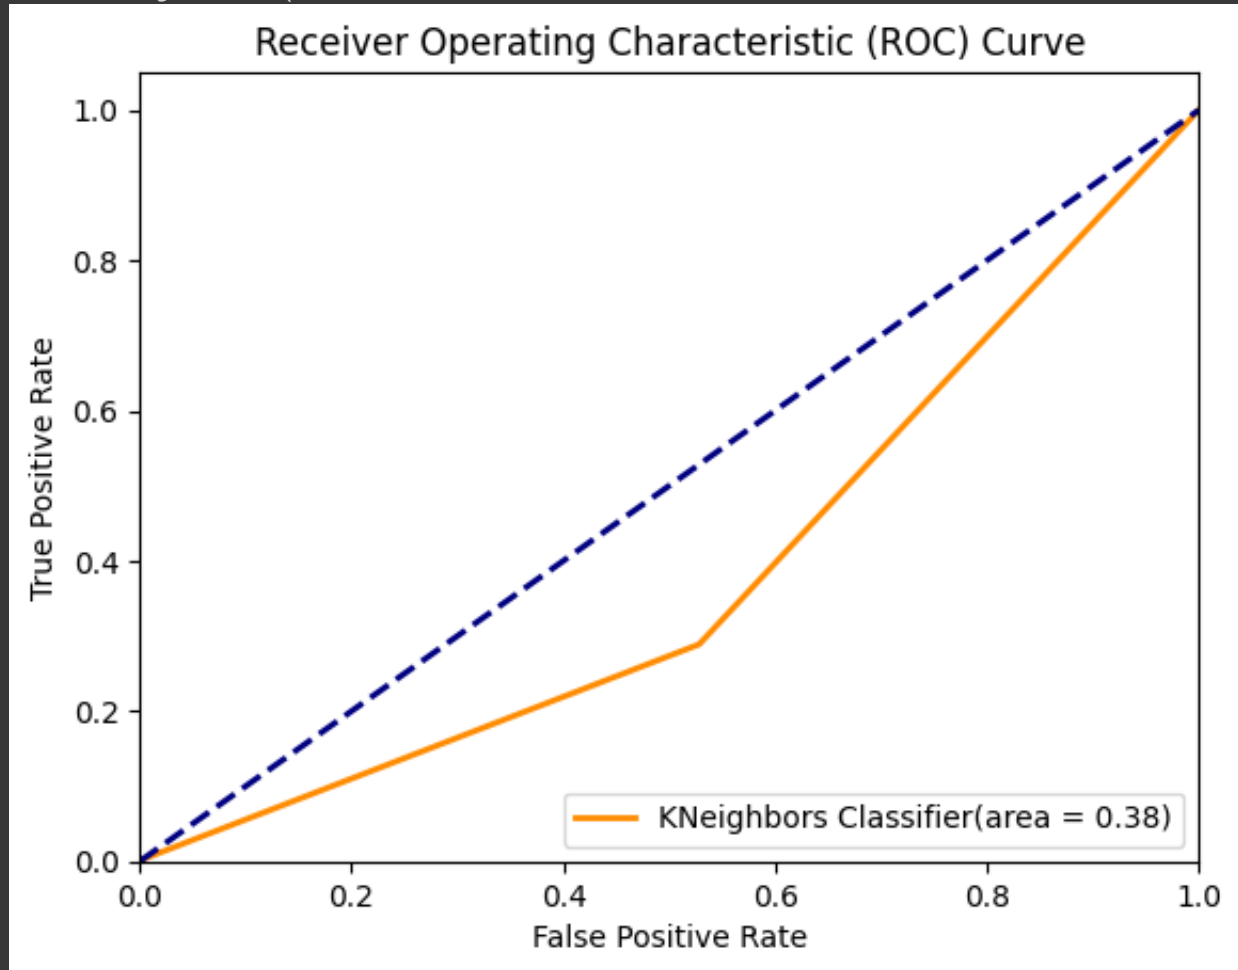
```
from sklearn.metrics import roc_curve, auc
import matplotlib.pyplot as plt

# Assuming knn_model is your classifier and X_test, y_test are your test data
y_score = knn_model.predict_proba(X_test)[:,1]
fpr, tpr, _ = roc_curve(y_test, y_score)
roc_auc = auc(fpr, tpr)

plt.figure()
plt.plot(fpr, tpr, color='darkorange', lw=2, label='KNeighbors Classifier(area = %0.2f)' % roc_auc)
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend(loc="lower right")
plt.show()
```



```
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:486: UserWarning: X h  
warnings.warn(
```



```
# Importing the RandomForestClassifier from sklearn.ensemble module
from sklearn.ensemble import RandomForestClassifier

# Creating a RandomForestClassifier model with specified parameters
random_forest_model = RandomForestClassifier(n_estimators=100, max_leaf_nodes=18)

# Training the RandomForestClassifier model on the normalized training data
random_forest_model.fit(x_train_normalized, y_train)

# Predicting the target labels for the normalized test data using the trained model
random_forest_model.predict(x_test_normalized)

# Calculating the accuracy score of the RandomForestClassifier model on the normalized test data
random_forest_model.score(x_test_normalized, y_test)
```

```
➔ /usr/local/lib/python3.10/dist-packages/sklearn/base.py:1474: DataConversionWarning:
    return fit_method(estimator, *args, **kwargs)
0.7421052631578947
```

```
import joblib
model_file_path = "/content/drive/My Drive/random_forest_model.pkl"
joblib.dump(random_forest_model, model_file_path)
```

```
➔ ['/content/drive/My Drive/random_forest_model.pkl']
```

```
# Random Forest – performance on the test data
from sklearn.metrics import accuracy_score, precision_score, recall_score, roc_auc_score
random_forest_pred = random_forest_model.predict(x_test_normalized)
random_forest_acc = accuracy_score(y_test, random_forest_pred )
random_forest_prec = precision_score(y_test, random_forest_pred )
random_forest_recall = recall_score(y_test, random_forest_pred )
random_forest_roc = roc_auc_score(y_test, random_forest_pred )
random_forest_f1 = f1_score(y_test, random_forest_pred )
print(random_forest_acc)
print(random_forest_prec)
print(random_forest_recall)
print(random_forest_roc)
print(random_forest_f1)
```

```
0.7140350877192982
0.6461794019933554
0.7749003984063745
0.7205222995166669
0.7047101449275363
```


```
from sklearn.metrics import roc_curve, auc
import matplotlib.pyplot as plt

# Assuming random_forest_model is your classifier and X_test, y_test are your test data
y_score = random_forest_model.predict_proba(X_test)[:, 1]
fpr, tpr, _ = roc_curve(y_test, y_score)
roc_auc = auc(fpr, tpr)

plt.figure()
plt.plot(fpr, tpr, color='darkorange', lw=2, label='RandomForestClassifier(area = %0.2f)' % roc_auc)
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend(loc="lower right")
plt.show()
```




```
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:486: UserWarning: X h
warnings.warn(
```



```
# Softmax Regression classifier
from sklearn.linear_model import LogisticRegression

softmax_regression = LogisticRegression(multi_class='multinomial', solver='lbfgs')
softmax_regression.fit(x_train_normalized, y_train)

# Get the predicted labels for training instances
y_prediction_labels = softmax_regression.predict(x_train_normalized)
```



```
/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:1300: Dat
y = column_or_1d(y, warn=True)
```

```
y_prediction_labels
```

```
⇒ array([0., 1., 0., ..., 0., 1., 0.]
```

```
# Get the predicted probabilities of both classes for training instances
y_prediction_probs = softmax_regression.predict_proba(x_train_normalized)
y_prediction_probs
```

```
⇒ array([[0.91584586, 0.08415414],
        [0.25960261, 0.74039739],
        [0.63893618, 0.36106382],
        ...,
        [0.84732268, 0.15267732],
        [0.40201416, 0.59798584],
        [0.71385464, 0.28614536]])
```

```
# Accuracy
import sklearn
sklearn.metrics.accuracy_score(y_train,y_prediction_labels)
```

```
⇒ 0.7267543859649123
```

```
# Build multi-layer neural network for binary classification
# Practice 2.1: Load training data
import pandas as pd
from sklearn.utils import shuffle

train_data = X_train
train_data.columns = ['Label'] + [f'F{i}' for i in range(1, 12)]
train_data = shuffle(train_data) # Shuffle data
train_data
```



| | Label | F1 | F2 | F3 | F4 | F5 | F6 | F7 | F8 | F9 |
|------|-------|-----|-----|------------|------|----------|----------|----------|----------|----------|
| 106 | 1 | 19 | 7 | 172.720000 | 72.6 | 0.187024 | 0.000000 | 0.000000 | 2.500000 | 1.800000 |
| 3544 | 1 | 6 | 7 | 179.000000 | 80.0 | 0.000000 | 0.000000 | 0.000000 | 2.538324 | 3.020000 |
| 2029 | 1 | 19 | 5 | 166.878000 | 79.3 | 0.187024 | 0.146335 | 0.066236 | 2.538324 | 2.364438 |
| 142 | 0 | 19 | 7 | 157.480000 | 49.9 | 0.187024 | 0.000000 | 0.000000 | 2.500000 | 2.400000 |
| 260 | 1 | 19 | 6 | 187.960000 | 83.5 | 0.187024 | 0.000000 | 0.000000 | 2.500000 | 2.000000 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 4380 | 0 | 6 | 6 | 168.047778 | 97.0 | 0.187024 | 0.146335 | 0.066236 | 2.538324 | 1.950000 |
| 5222 | 1 | 19 | 6 | 180.009800 | 71.0 | 0.000000 | 0.000000 | 0.000000 | 2.538324 | 2.364438 |
| 1738 | 1 | 19 | 6 | 173.736000 | 84.1 | 1.000000 | 1.000000 | 0.000000 | 2.500000 | 2.800000 |
| 359 | 0 | 18 | 4 | 170.180000 | 67.1 | 0.187024 | 0.000000 | 0.000000 | 3.000000 | 2.200000 |
| 2935 | 0 | 13 | 5 | 168.047778 | 56.0 | 0.000000 | 0.000000 | 0.000000 | 2.538324 | 1.720000 |

4560 rows × 12 columns

```
import pandas as pd

test_data = x_test_normalized
test_data.columns = ['Label'] + [f'F{i}' for i in range(1, 12)]
test_data
```



| | Label | F1 | F2 | F3 | F4 | F5 | F6 | F7 | F8 |
|------|-------|--------|----------|----------|----------|----------|----------|----------|----------|
| 0 | 0.0 | 0.4375 | 0.142857 | 0.332917 | 0.241880 | 0.000000 | 0.000000 | 0.000000 | 0.562871 |
| 1 | 1.0 | 0.2500 | 0.714286 | 0.833333 | 0.339254 | 0.000000 | 0.000000 | 0.000000 | 0.562871 |
| 2 | 0.0 | 0.3125 | 0.571429 | 0.465023 | 0.297167 | 0.000000 | 0.146335 | 1.000000 | 0.562871 |
| 3 | 1.0 | 1.0000 | 0.285714 | 0.546250 | 0.290256 | 0.000000 | 0.000000 | 1.000000 | 0.562871 |
| 4 | 0.0 | 1.0000 | 0.857143 | 0.300000 | 0.345543 | 0.000000 | 0.000000 | 0.000000 | 0.562871 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 1135 | 1.0 | 0.6875 | 0.285714 | 0.366667 | 0.172771 | 0.187024 | 0.146335 | 0.066236 | 0.562871 |
| 1136 | 1.0 | 0.3125 | 0.142857 | 0.465023 | 0.200415 | 0.000000 | 0.146335 | 0.000000 | 0.562871 |
| 1137 | 1.0 | 1.0000 | 0.285714 | 0.637500 | 0.373186 | 0.000000 | 1.000000 | 0.000000 | 0.545455 |
| 1138 | 1.0 | 0.6875 | 0.571429 | 0.445833 | 0.124395 | 0.187024 | 0.146335 | 0.066236 | 0.562871 |
| 1139 | 1.0 | 0.6875 | 0.571429 | 0.662500 | 0.262612 | 0.187024 | 0.146335 | 0.066236 | 0.562871 |

1140 rows × 10 columns

```
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense

model = Sequential()
model.add(Dense(100, input_shape=(22,), activation='relu'))

model.summary()
```

➞ Model: "sequential_3"

| Layer (type) | Output Shape | Param # |
|-------------------------------------|--------------|---------|
| dense_7 (Dense) | (None, 100) | 2300 |
| Total params: 2300 (8.98 KB) | | |
| Trainable params: 2300 (8.98 KB) | | |
| Non-trainable params: 0 (0.00 Byte) | | |

```

from keras.models import Sequential
from keras.layers import Dense

model = Sequential([
    Dense(100, input_shape=(22,), activation='relu'),
    Dense(100, activation='relu'),
    Dense(1, activation='sigmoid')
])

model.summary()

```

➞ Model: "sequential_4"

| Layer (type) | Output Shape | Param # |
|-------------------------------------|--------------|---------|
| dense_8 (Dense) | (None, 100) | 2300 |
| dense_9 (Dense) | (None, 100) | 10100 |
| dense_10 (Dense) | (None, 1) | 101 |
| Total params: 12501 (48.83 KB) | | |
| Trainable params: 12501 (48.83 KB) | | |
| Non-trainable params: 0 (0.00 Byte) | | |

```
# Compile the model for backpropagation, and start training model on data
from keras.models import Sequential
from keras.layers import Dense

model = Sequential([
    Dense(100, input_shape=(12,), activation='relu'),
    Dense(100, activation='relu'),
    Dense(1, activation='sigmoid')
])
# Compile model, use binary crossentropy, and stochastic gradient descent for opt
model.compile(loss='binary_crossentropy', optimizer='SGD', metrics=['accuracy'])

model.summary()
```

⇒ Model: "sequential_5"

| Layer (type) | Output Shape | Param # |
|-------------------------------------|--------------|---------|
| dense_11 (Dense) | (None, 100) | 1300 |
| dense_12 (Dense) | (None, 100) | 10100 |
| dense_13 (Dense) | (None, 1) | 101 |
| Total params: 11501 (44.93 KB) | | |
| Trainable params: 11501 (44.93 KB) | | |
| Non-trainable params: 0 (0.00 Byte) | | |

```
model.fit(X_train, y_train, epochs = 50)
```

⇒ Epoch 1/50
 143/143 [=====] - 5s 7ms/step - loss: 1.1637 - accur
 Epoch 2/50
 143/143 [=====] - 1s 6ms/step - loss: 0.6468 - accur
 Epoch 3/50
 143/143 [=====] - 1s 10ms/step - loss: 0.6447 - accu
 Epoch 4/50
 143/143 [=====] - 1s 7ms/step - loss: 0.6422 - accur
 Epoch 5/50
 143/143 [=====] - 1s 9ms/step - loss: 0.6389 - accur
 Epoch 6/50
 143/143 [=====] - 1s 7ms/step - loss: 0.6368 - accur
 Epoch 7/50
 143/143 [=====] - 1s 5ms/step - loss: 0.6357 - accur

```

Epoch 8/50
143/143 [=====] - 1s 4ms/step - loss: 0.6337 - accur
Epoch 9/50
143/143 [=====] - 2s 11ms/step - loss: 0.6350 - accur
Epoch 10/50
143/143 [=====] - 1s 10ms/step - loss: 0.6352 - accur
Epoch 11/50
143/143 [=====] - 2s 11ms/step - loss: 0.6318 - accur
Epoch 12/50
143/143 [=====] - 1s 7ms/step - loss: 0.6303 - accur
Epoch 13/50
143/143 [=====] - 1s 9ms/step - loss: 0.6317 - accur
Epoch 14/50
143/143 [=====] - 1s 9ms/step - loss: 0.6301 - accur
Epoch 15/50
143/143 [=====] - 1s 9ms/step - loss: 0.6287 - accur
Epoch 16/50
143/143 [=====] - 1s 5ms/step - loss: 0.6270 - accur
Epoch 17/50
143/143 [=====] - 1s 6ms/step - loss: 0.6241 - accur
Epoch 18/50
143/143 [=====] - 1s 6ms/step - loss: 0.6228 - accur
Epoch 19/50
143/143 [=====] - 1s 7ms/step - loss: 0.6219 - accur
Epoch 20/50
143/143 [=====] - 1s 6ms/step - loss: 0.6259 - accur
Epoch 21/50
143/143 [=====] - 1s 5ms/step - loss: 0.6224 - accur
Epoch 22/50
143/143 [=====] - 1s 4ms/step - loss: 0.6215 - accur
Epoch 23/50
143/143 [=====] - 0s 2ms/step - loss: 0.6243 - accur
Epoch 24/50
143/143 [=====] - 0s 2ms/step - loss: 0.6241 - accur
Epoch 25/50
143/143 [=====] - 0s 2ms/step - loss: 0.6183 - accur
Epoch 26/50
143/143 [=====] - 0s 2ms/step - loss: 0.6215 - accur
Epoch 27/50
143/143 [=====] - 0s 2ms/step - loss: 0.6167 - accur
Epoch 28/50
143/143 [=====] - 0s 2ms/step - loss: 0.6211 - accur
Epoch 29/50
143/143 [=====] - 0s 2ms/step - loss: 0.6108 - accur
Epoch 30/50
143/143 [=====] - 0s 2ms/step - loss: 0.6177 - accur

```



```
from tensorflow.keras.models import load_model
model.save("/content/drive/My Drive/tensor_keras.pkl")
#model_loaded = load_model("my_model.h5")
```

```
model.evaluate(X_train, y_train)
```

```
⇒ 143/143 [=====] - 0s 2ms/step - loss: 0.5957 - accur
[0.5956673622131348, 0.6780701875686646]
```

```
train_predictions = model.predict(X_train)
test_predictions = model.predict(X_test)
print("train_predictions: ", X_train) # Print train_predictions
```

```
⇒ 143/143 [=====] - 0s 2ms/step
36/36 [=====] - 0s 2ms/step
train_predictions:      Label  F1  F2      F3      F4      F5
5437      0  19   6  160.020000  60.91  0.000000  0.000000  0.000000
2979      1  13   1  178.000000  82.00  0.000000  0.000000  0.000000
4743      0  14   5  168.047778  72.00  0.187024  0.146335  0.066236
2668      1  19   6  175.260000  85.30  0.000000  0.000000  0.000000
3264      1   6   7  168.047778  59.00  0.000000  0.146335  0.000000
...      ... ..
1180      1  12   3  170.942000  79.00  0.187024  0.146335  0.066236
3441      0   2   5  170.180000 128.60  1.000000  1.000000  0.000000
1344      0  12   5  153.924000  60.00  0.187024  0.146335  0.066236
4623      0  19   3  160.020000  60.00  0.000000  0.000000  0.000000
1289      0  12   5  157.988000  60.00  0.187024  0.146335  0.066236

      F8      F9  F10  F11
5437  2.538324  2.200000   5   0
2979  2.538324  2.470000   0   0
4743  2.538324  2.120000   0   0
2668  2.538324  2.500000   5   0
3264  2.538324  2.364438   4   2
...      ... ..
1180  2.538324  2.364438   0   0
3441  2.500000  2.700000   0   2
1344  2.538324  1.370000   0   0
4623  2.538324  2.800000   5   2
1289  2.538324  2.990000   0   0
```

```
[4560 rows x 12 columns]
```

```
import numpy as np

# Convert predicted probabilities to binary labels using a threshold of 0.5
train_prediction_labels = (train_predictions > 0.5).astype(int)
test_prediction_labels = (test_predictions > 0.5).astype(int)
```

```
from sklearn.metrics import accuracy_score

# Calculate and print the training accuracy
print("Training accuracy: ", accuracy_score(y_train, train_prediction_labels))

# Calculate and print the test accuracy
print("Test accuracy: ", accuracy_score(y_test, test_prediction_labels))
```

➞ Training accuracy: 0.6780701754385965
Test accuracy: 0.6912280701754386

```
from sklearn.decomposition import PCA

# Instantiate PCA with desired number of components
pca = PCA(n_components=2)


# Fit PCA to training data and transform it in one step
X_reduced = pca.fit_transform(X_train)
```

```
# Initial Shape
X_train.shape
print("Initial shape:", X_train.shape)
# Reduced Shape
X_reduced.shape
print("Reduce shape:", X_reduced.shape)
```

➞ Initial shape: (4560, 12)
Reduce shape: (4560, 2)

```
from sklearn.discriminant_analysis import LinearDiscriminantAnalysis

# Instantiate and train LDA model
lda = LinearDiscriminantAnalysis()
lda.fit(X_reduced, y_train)
```

 /usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:1300: Data y = column or 1d(y, warn=True)


▼ LinearDiscriminantAnalysis ⓘ ?

LinearDiscriminantAnalysis()

```
import joblib
model_file_path = "/content/drive/My Drive/LinearDiscriminantAnalysis_model.pkl"
joblib.dump(knn_model, model_file_path)
```

 ['/content/drive/My Drive/LinearDiscriminantAnalysis_model.pkl']

```
# Score
lda.score(X_reduced, y_train)
```

 0.6515350877192982

```
# Evaluation Metrics
TabularData = {'Methods': ['Decision Tree', 'SVM', 'KNN', 'Random Forest', 'Logis
'Accuracy': [decision_tree_model_acc, svm_acc, knn_acc, random_fore
'Precision': [decision_tree_model_prec, svm_prec, knn_prec, random_
'Recall': [decision_tree_model_recall, svm_recall, knn_recall, rand
'F1-score': [decision_tree_model_f1, svm_f1, knn_f1, random_forest_
'AUC score': [decision_tree_model_roc, svm_roc, knn_roc, random_for

evaluation_metrics = pd.DataFrame(TabularData)
```

evaluation_metrics



| | Methods | Accuracy | Precision | Recall | F1-score | AUC score |
|---|---------------------|----------|-----------|----------|----------|-----------|
| 0 | Decision Tree | 0.637719 | 0.560217 | 0.824701 | 0.667204 | 0.657648 |
| 1 | SVM | 0.718421 | 0.661319 | 0.739044 | 0.698024 | 0.720619 |
| 2 | KNN | 0.714912 | 0.665421 | 0.709163 | 0.686596 | 0.714300 |
| 3 | Random Forest | 0.714035 | 0.646179 | 0.774900 | 0.704710 | 0.720522 |
| 4 | Logistic regression | 0.703509 | 0.623867 | 0.822709 | 0.709622 | 0.716214 |

```
# Cross Validation for Decision Tree
```

```
from sklearn.tree import DecisionTreeClassifier
```

```
from sklearn.model_selection import cross_val_score
```

```
CV_scores_decision_tree_model = cross_val_score(estimator = decision_tree_model,  
print("CV_scores: ", CV_scores_decision_tree_model)
```



```
CV_scores: [0.67763158 0.67763158 0.67105263 0.65131579 0.68421053 0.7565789  
0.72368421 0.63157895 0.65131579 0.80921053 0.73684211 0.69736842  
0.67763158 0.71710526 0.72368421 0.65789474 0.68421053 0.625  
0.79605263 0.68421053 0.73026316 0.69078947 0.66447368 0.73026316  
0.65789474 0.72368421 0.69736842 0.63815789 0.71710526 0.69736842]
```

```
# Predicting on test data using the decision tree model
```

```
y_test_pred_dt = decision_tree_model.predict(x_test_normalized)
```



```
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:486: UserWarning: X h  
warnings.warn(
```

```
from sklearn import metrics
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score

# Calculating accuracy
decision_tree_acc_cv = accuracy_score(y_test_pred_dt, y_test)

# Calculating precision
decision_tree_prec_cv = precision_score(y_test_pred_dt, y_test)

# Calculating recall
decision_tree_recall_cv = recall_score(y_test_pred_dt, y_test)

# Calculating F1 score
decision_tree_f1_cv = f1_score(y_test_pred_dt, y_test)

# Calculating ROC AUC score (assuming you have imported roc_auc_score)
from sklearn.metrics import roc_auc_score
decision_tree_roc_cv = roc_auc_score(y_test_pred_dt, y_test)

# Printing the calculated metrics
print("Accuracy: ", decision_tree_acc_cv)
print("Precision:", decision_tree_prec_cv)
print("Recall:", decision_tree_recall_cv)
print("F1 Score:", decision_tree_f1_cv)
print("ROC AUC Score:", decision_tree_roc_cv)
```

⇒ Accuracy: 0.637719298245614
Precision: 0.8247011952191236
Recall: 0.5602165087956699
F1 Score: 0.6672038678485093
ROC AUC Score: 0.6703825686122988


```
# Plotting the boxplot of cross-validation scores
plt.boxplot(CV_scores_decision_tree_model)

# Adding title to the plot
plt.title("10-fold cross validation accuracy")

# Adding label to x-axis
plt.xlabel("Decision Tree")

# Adding label to y-axis
plt.ylabel("Accuracy")

# Displaying the plot
plt.show()
```



```
# Cross validation for Logistic regression
from sklearn.model_selection import cross_val_score
from sklearn.linear_model import LogisticRegression
```

```
# Creating an instance of Logistic Regression model
logistic_model = LogisticRegression()

# Performing cross-validation
CV_scores_lr = cross_val_score(logistic_model, X_train, y_train, cv=10, scoring='

# Printing the cross-validation scores
print("CV_scores:", CV_scores_lr)
```

```
➞ /usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:1300: Dat
y = column_or_1d(y, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py:469
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regres

```
n_iter_i = _check_optimize_result(
/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:1300: Dat
y = column_or_1d(y, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py:469
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regres

```
n_iter_i = _check_optimize_result(
/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:1300: Dat
y = column_or_1d(y, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py:469
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regres

```
n_iter_i = _check_optimize_result(
/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:1300: Dat
y = column_or_1d(y, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py:469
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regres

```

n_iter_i = _check_optimize_result(
/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:1300: Dat
y = column_or_1d(y, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py:469
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regres

```

n_iter_i = _check_optimize_result(
/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:1300: Dat
y = column_or_1d(y, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py:469
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regres

```

# Accuracy, Precision, Recall and F1-score
# Predicting on test data using the random forest model
y_test_pred3 = random_forest_model.predict(x_test_normalized)

# Calculating and printing accuracy
print("Accuracy: ", accuracy_score(y_test_pred3, y_test))

# Calculating and printing precision
print("Precision:", precision_score(y_test_pred3.astype(int), y_test.astype(int))

# Calculating and printing recall
print("Recall:", recall_score(y_test_pred3.astype(int), y_test.astype(int)))

# Calculating and printing F1-score
print("F1 Score:", f1_score(y_test_pred3.astype(int), y_test.astype(int)))

```

```

➡ Accuracy: 0.7140350877192982
Precision: 0.7749003984063745
Recall: 0.6461794019933554
F1 Score: 0.7047101449275363
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:486: UserWarning: X h
warnings.warn(

```



```
# Predicting on the test data using the random forest model
y_test_pred_lr = random_forest_model.predict(X_test)
```

➞ /usr/local/lib/python3.10/dist-packages/sklearn/base.py:486: UserWarning: X has feature names, but LogisticRegression was fitted without feature names

```
from sklearn import metrics
from sklearn.metrics import accuracy_score
from sklearn.metrics import precision_score, recall_score, f1_score
# Calculating accuracy
logistic_model_acc_cv = metrics.accuracy_score(y_test_pred_lr, y_test)
# Calculating precision
logistic_model_prec_cv = precision_score(y_test_pred_lr, y_test)

# Calculating recall
logistic_model_recall_cv = recall_score(y_test_pred_lr, y_test)

# Calculating F1 score
logistic_model_f1_cv = f1_score(y_test_pred_lr, y_test)

# Calculating ROC AUC score
logistic_model_roc_cv = roc_auc_score(y_test_pred_lr, y_test)

# Printing the calculated metrics
print("Accuracy: ", logistic_model_acc_cv)
print("Precision:", logistic_model_prec_cv)
print("Recall:", logistic_model_recall_cv)
print("F1 Score:", logistic_model_f1_cv)
print("ROC AUC Score:", logistic_model_roc_cv)
```

➞ Accuracy: 0.4850877192982456
Precision: 0.9203187250996016
Recall: 0.4578790882061447
F1 Score: 0.6115155526141628
ROC AUC Score: 0.5762677883778815


```
# Plotting the boxplot of cross-validation scores for logistic regression
plt.boxplot(CV_scores_lr)

# Adding title to the plot
plt.title("10-fold cross validation accuracy")

# Adding label to x-axis
plt.xlabel("Logistic Regression")

# Adding label to y-axis
plt.ylabel("Accuracy")

# Displaying the plot
plt.show()
```



```
# Cross validation for Random Forest
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import cross_val_score

model1 = random_forest_model
CV_scores_random_forest_model = cross_val_score(estimator = model1, X = X_train,
print("CV_scores: ", CV_scores_random_forest_model)
```

```
➞ /usr/local/lib/python3.10/dist-packages/sklearn/base.py:1474: DataConversionW
    return fit_method(estimator, *args, **kwargs)
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:1474: DataConversionW
    return fit_method(estimator, *args, **kwargs)
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:1474: DataConversionW
    return fit_method(estimator, *args, **kwargs)
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:1474: DataConversionW
    return fit_method(estimator, *args, **kwargs)
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:1474: DataConversionW
    return fit_method(estimator, *args, **kwargs)
CV_scores: [0.71820175 0.72807018 0.72697368 0.73135965 0.7247807 ]
```

```
# Predicting on the test data using the random forest model
y_test_pred_rf = random_forest_model.predict(X_test)
```

```
➞ /usr/local/lib/python3.10/dist-packages/sklearn/base.py:486: UserWarning: X h
    warnings.warn(
```

```
from sklearn import metrics
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score

# Calculating accuracy
random_forest_acc_cv = accuracy_score(y_test_pred_rf, y_test)

# Calculating precision
random_forest_prec_cv = precision_score(y_test_pred_rf, y_test)

# Calculating recall
random_forest_recall_cv = recall_score(y_test_pred_rf, y_test)

# Calculating F1 score
random_forest_f1_cv = f1_score(y_test_pred_rf, y_test)

# Calculating ROC AUC score (assuming you have imported roc_auc_score)
random_forest_roc_cv = roc_auc_score(y_test_pred_rf, y_test)


# Printing the calculated metrics
print("Accuracy: ", random_forest_acc_cv)
print("Precision:", random_forest_prec_cv)
print("Recall:", random_forest_recall_cv)
print("F1 Score:", random_forest_f1_cv)
print("ROC AUC Score:", random_forest_roc_cv)
```

⇒ Accuracy: 0.4850877192982456
Precision: 0.9203187250996016
Recall: 0.4578790882061447
F1 Score: 0.6115155526141628
ROC AUC Score: 0.5762677883778815

```
plt.boxplot(CV_scores_random_forest_model)
plt.title("10-fold cross validation accuracy")
plt.xlabel("Random Forest")
plt.ylabel("Accuracy")
```



```
Text(0, 0.5, 'Accuracy')
```



```
# Evaluation Metrics
```

```
TabularData_cv = {'Methods': ['Decision Tree', 'Random Forest', 'Logistic Model'],
                  'Accuracy': [decision_tree_acc_cv, random_forest_acc_cv, logistic_model_a
                  'Precision': [decision_tree_prec_cv, random_forest_prec_cv, logistic_model
                  'Recall': [decision_tree_recall_cv, random_forest_recall_cv, logistic_mode
                  'F1-score': [decision_tree_f1_cv, random_forest_f1_cv, logistic_model_f1_c
                  'AUC score': [decision_tree_roc_cv, random_forest_roc_cv, logistic_model_r
```

```
evaluation_metrics_cv =pd.DataFrame(TabularData_cv)
```

evaluation_metrics_cv



| | Methods | Accuracy | Precision | Recall | F1-score | AUC score |
|---|----------------|----------|-----------|----------|----------|-----------|
| 0 | Decision Tree | 0.637719 | 0.824701 | 0.560217 | 0.667204 | 0.670383 |
| 1 | Random Forest | 0.485088 | 0.920319 | 0.457879 | 0.611516 | 0.576268 |
| 2 | Logistic Model | 0.485088 | 0.920319 | 0.457879 | 0.611516 | 0.576268 |

```
pip install gradio
```



```
Requirement already satisfied: gradio in /usr/local/lib/python3.10/dist-packa
Requirement already satisfied: aiofiles<24.0,>=22.0 in /usr/local/lib/python3
Requirement already satisfied: altair<6.0,>=4.2.0 in /usr/local/lib/python3.1
Requirement already satisfied: fastapi in /usr/local/lib/python3.10/dist-pack
Requirement already satisfied: ffmpeg in /usr/local/lib/python3.10/dist-packag
Requirement already satisfied: gradio-client==0.16.3 in /usr/local/lib/python
Requirement already satisfied: httpx>=0.24.1 in /usr/local/lib/python3.10/dis
Requirement already satisfied: huggingface-hub>=0.19.3 in /usr/local/lib/pyth
Requirement already satisfied: importlib-resources<7.0,>=1.3 in /usr/local/li
Requirement already satisfied: jinja2<4.0 in /usr/local/lib/python3.10/dist-p
Requirement already satisfied: markupsafe~=2.0 in /usr/local/lib/python3.10/d
Requirement already satisfied: matplotlib~=3.0 in /usr/local/lib/python3.10/d
Requirement already satisfied: numpy~=1.0 in /usr/local/lib/python3.10/dist-p
Requirement already satisfied: orjson~=3.0 in /usr/local/lib/python3.10/dist-
Requirement already satisfied: packaging in /usr/local/lib/python3.10/dist-pa
Requirement already satisfied: pandas<3.0,>=1.0 in /usr/local/lib/python3.10/
Requirement already satisfied: pillow<11.0,>=8.0 in /usr/local/lib/python3.10
Requirement already satisfied: pydantic>=2.0 in /usr/local/lib/python3.10/dis
Requirement already satisfied: pydub in /usr/local/lib/python3.10/dist-packag
Requirement already satisfied: python-multipart>=0.0.9 in /usr/local/lib/pyth
Requirement already satisfied: pyyaml<7.0,>=5.0 in /usr/local/lib/python3.10/
Requirement already satisfied: ruff>=0.2.2 in /usr/local/lib/python3.10/dist-
Requirement already satisfied: semantic-version~=2.0 in /usr/local/lib/python
Requirement already satisfied: tomlkit==0.12.0 in /usr/local/lib/python3.10/d
Requirement already satisfied: typer<1.0,>=0.12 in /usr/local/lib/python3.10/
Requirement already satisfied: typing-extensions~=4.0 in /usr/local/lib/pytho
Requirement already satisfied: urllib3~=2.0 in /usr/local/lib/python3.10/dist
Requirement already satisfied: uvicorn>=0.14.0 in /usr/local/lib/python3.10/d
Requirement already satisfied: fsspec in /usr/local/lib/python3.10/dist-packa
Requirement already satisfied: websockets<12.0,>=10.0 in /usr/local/lib/pytho
Requirement already satisfied: entrypoints in /usr/local/lib/python3.10/dist-
Requirement already satisfied: jsonschema>=3.0 in /usr/local/lib/python3.10/d
Requirement already satisfied: toolz in /usr/local/lib/python3.10/dist-packag
Requirement already satisfied: anyio in /usr/local/lib/python3.10/dist-packag
Requirement already satisfied: certifi in /usr/local/lib/python3.10/dist-pack
```

Requirement already satisfied: httpcore==1.* in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: idna in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: sniffio in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: h11<0.15,>=0.13 in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: filelock in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: requests in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: tqdm>=4.42.1 in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: kiwisolver>=1.0.1 in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: python-dateutil>=2.7 in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: tzdata>=2022.1 in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: annotated-types>=0.4.0 in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: pydantic-core==2.18.2 in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: click>=8.0.0 in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: shellingham>=1.3.0 in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: rich>=10.11.0 in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: starlette<0.38.0,>=0.37.2 in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: fastapi-cli>=0.0.2 in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: ujson!=4.0.2,!=4.1.0,!=4.2.0,!=4.3.0,!=5.0.0,!=5.1.0 in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: email-validator>=2.0.0 in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: dnspython>=2.0.0 in /usr/local/lib/python3.10/dist-packages

